

Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, Peter J. Liu

Presentation by: Rafael Ferreira

Date: 14/05/2020

Content

1. T5 – Text-to-Text Transfer Transformer
2. The Transformer
3. T5 Applications
4. Tasks and Datasets
5. T5 Model - Baseline
6. Architectures
7. Unsupervised Objectives
8. Pre-training Datasets
9. Training strategies
10. Scaling
11. Conclusion
12. Research Opportunities

T5 – **T**ext-to-**T**ext **T**ransfer **T**ransformer



Text-To-Text – Relates to the input and output formats of the model.
The model receives input to utilize as context to then produce an output.



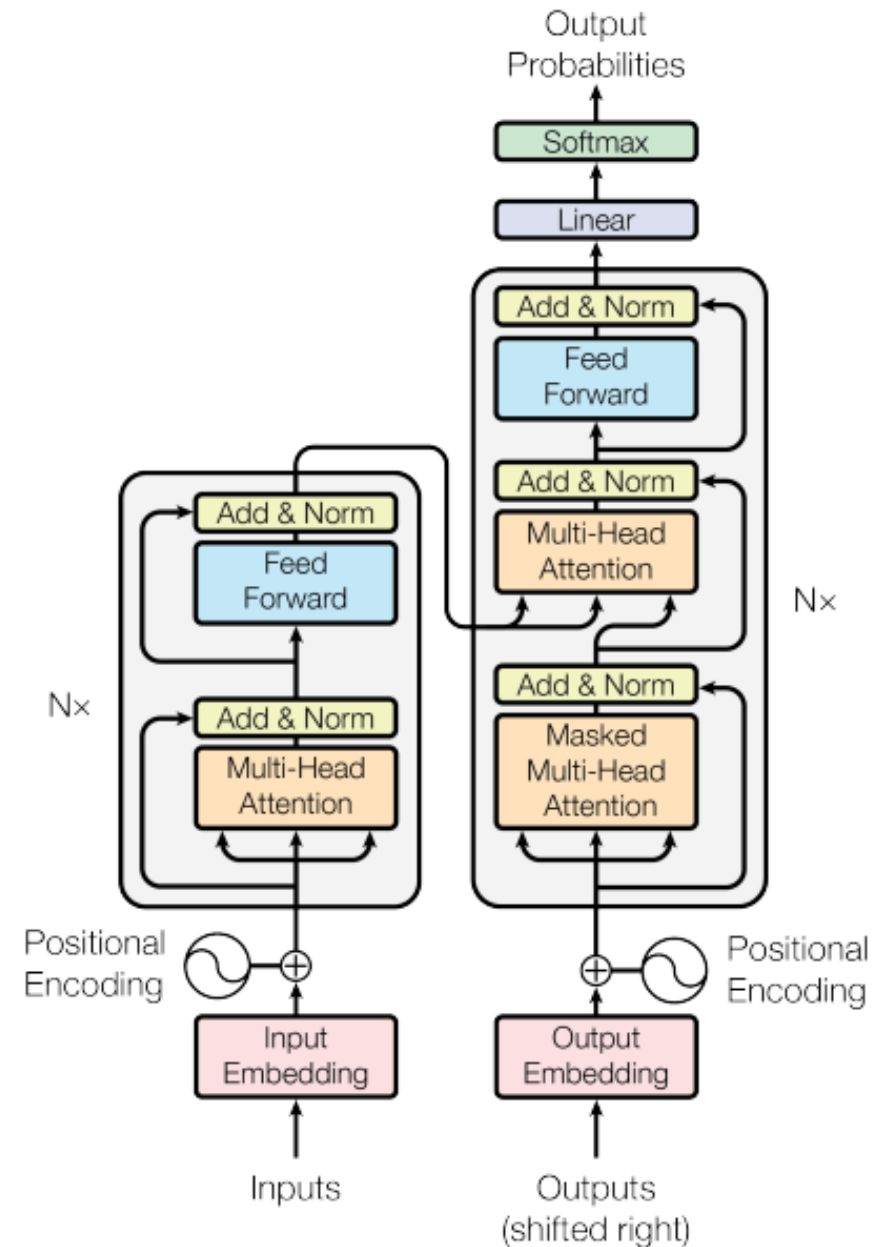
Transfer – Transfer learning. A Transfer learning model is a model that is first pre-trained on large amounts of data (unsupervised) and is then fine-tuned on a task.



Transformer – The architecture of the model.

The Transformer

- The encoder-decoder version of T5 follows the architecture proposed in **Attention is all you need**¹.
- One of the most important parts of this architecture is the **self-attention** mechanism. A variant of attention that replaces each element in a sequence by a weighted average of the rest of the sequence.
- A tutorial on the attention mechanism and the transformer architecture can be seen [here](#).



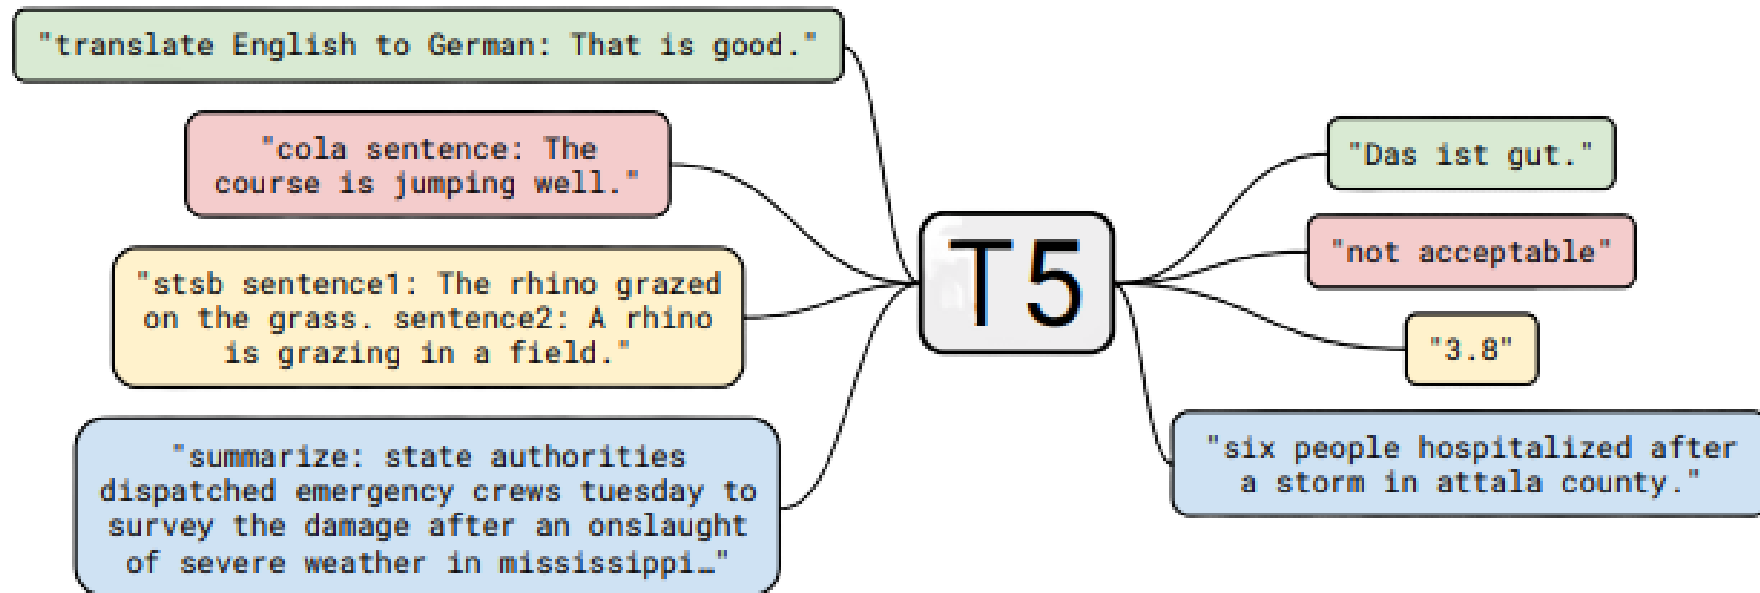
¹Ashish Vaswani et al., "Attention Is All You Need"

T5 Applications

One of the advantages of a text-to-text format is that it provides a **consistent training objective** both for pre-training and fine-tuning.

T5 is also able to handle **multiple tasks** by adding task-specific prefixes to the original input sequence.

- Machine Translation
- “Regression”
- Summarization
- Classification
- Question Answering
- Coreference Resolution



Tasks and Datasets

- **GLUE**¹ and **SuperGLUE**² – collection of text classification tasks.
- **CNNDM**³ (CNN/Daily Mail) – summarization task.
- **SQuAD**⁴ – extractive question-answering dataset.
- **EnDe**, **EnFr** and **EnRo** – translation tasks from English to German (Deutsch), French, and Romanian respectively.

¹Wang et al., “GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding”

²Wang et al., “SuperGLUE: A Stickier Benchmark for General-Purpose Language Understanding Systems”

³Hermann et al., “Teaching Machines to Read and Comprehend”

⁴Rajpurkar et al., “Squad: 100,000+ questions for machine comprehension of text”

T5 Model - Baseline


- **Architecture** - Encoder-decoder transformer architecture.
 - 24 layers / hidden dimension size 768 / 12 attention heads,
 - 220 Million parameters (similar to BERT_{BASE}).
- **Training** - Maximum likelihood.
 - Pre-trained model for 2^{19} steps and fine-tuned for 2^{18} steps,
 - Maximum sequence length of 512 tokens and batch size of 128 sequences.
- **Unsupervised Objective** - “denoising”, predicting missing or corrupted tokens from the input.

	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Baseline average	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Baseline standard deviation	0.235	0.065	0.343	0.416	0.112	0.090	0.108
No pre-training	66.22	17.60	50.31	53.04	25.86	39.77	24.04

Architectures

- **Attention Masks**

- Used to zero out certain weights to attend only to inputs it considers necessary at each step.



- **Model Structures:**

- **Encoder-decoder:**


- Encoder uses fully-visible attention mask
- Decoder uses causal attention mask

- **Language model:**

- Next Step Prediction

- **Prefix LM**

- Fully-visible attention mask on prefix
- Causal attention mask outside prefix



Architectures

- Considering the number of layers and parameters as \mathbf{L} and \mathbf{P} respectively of a BERT_{BASE} model, and \mathbf{M} as the number of FLOPs.

Architecture	Objective	Params	Cost	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Encoder-decoder	Denoising	$2P$	M	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Enc-dec, shared	Denoising	P	M	82.81	18.78	80.63	70.73	26.72	39.03	27.46
Enc-dec, 6 layers	Denoising	P	$M/2$	80.88	18.97	77.59	68.42	26.38	38.40	26.95
Language model	Denoising	P	M	74.70	17.93	61.14	55.02	25.09	35.28	25.86
Prefix LM	Denoising	P	M	81.82	18.61	78.94	68.11	26.43	37.98	27.39

Unsupervised Objectives

- Decision of which algorithm to use as the unsupervised objective during pre-training is essential, because it provides the method through which the model “**learns the knowledge**”.
- Examples of unsupervised objectives for the sentence:
 - “Thank you for inviting me to your party last week.”

Objective	Inputs	Targets
Prefix language modeling	Thank you for inviting	me to your party last week .
BERT-style	Thank you <M> <M> me to your party apple week .	<i>(original text)</i>
Deshuffling	party me for your to . last fun you inviting week Thank	<i>(original text)</i>
I.i.d. noise, mask tokens	Thank you <M> <M> me to your party <M> week .	<i>(original text)</i>
I.i.d. noise, replace spans	Thank you <X> me to your party <Y> week .	<X> for inviting <Y> last <Z>
I.i.d. noise, drop tokens	Thank you me to your party week .	for inviting last
Random spans	Thank you <X> to <Y> week .	<X> for inviting me <Y> your party last <Z>

Unsupervised Objectives

- Deshuffling objectives perform significantly worse.
- Denoising objectives are the best but comparable between themselves.
- The major advantage of span corruption comes from their **lower computational cost**, because on average, span corruption produces shorter sequences.

Objective	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
Prefix language modeling	80.69	18.94	77.99	65.27	26.86	39.73	27.49
BERT-style [Devlin et al., 2018]	82.96	19.17	80.65	69.85	26.78	40.03	27.41
Deshuffling	73.17	18.59	67.61	58.47	26.11	39.30	25.62
MASS-style [Song et al., 2019]	82.32	19.16	80.10	69.28	26.79	39.89	27.55
★ Replace corrupted spans	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Drop corrupted tokens	84.44	19.31	80.52	68.67	27.07	39.76	27.82

Pre-Training Datasets

- The pre-training dataset is another essential part of training, alongside the pre-training objective.
- There exists various large-scale unlabeled datasets. More data being produced each month.
- C4
- RealNews-like (C4 news)
- WebText-Like (C4 Reddit)
- Wikipedia
- Toronto Book Corpus

Dataset	Size	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ C4	745GB	83.28	19.24	80.88	71.36	26.98	39.82	27.65
C4, unfiltered	6.1TB	81.46	19.14	78.78	68.04	26.55	39.34	27.21
RealNews-like	35GB	83.83	19.23	80.39	72.38	26.75	39.90	27.48
WebText-like	17GB	84.03	19.31	81.42	71.40	26.80	39.74	27.59
Wikipedia	16GB	81.85	19.31	81.29	68.01	26.94	39.69	27.67
Wikipedia + TBC	20GB	83.65	19.28	82.08	73.24	26.77	39.63	27.57

- Main takeaway is that **pre-training on in-domain** unlabeled data can improve performance on downstream tasks.

Pre-Training Datasets

- **Influence of repeating data** - Tested using the C4 dataset truncated at different sizes.

Number of tokens	Repeats	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Full dataset	0	83.28	19.24	80.88	71.36	26.98	39.82	27.65
2^{29}	64	82.87	19.19	80.97	72.03	26.83	39.74	27.63
2^{27}	256	82.62	19.20	79.78	69.97	27.02	39.71	27.33
2^{25}	1,024	79.55	18.57	76.27	64.76	26.38	39.56	26.80
2^{23}	4,096	76.34	18.33	70.92	59.29	26.37	38.84	25.81

- **Performance generally degrades** as the **size** of the dataset **decreases**.
- When the dataset is repeated 64 times it surpasses the full dataset in some tasks showing that some amount of repetition might not be harmful.

Training strategies

- **Fine-tuning methods**

- **Adapter Layers**¹ – dense ReLU blocks added after each existing FFN in each block of the Transformer. Only the **adapter layers and layer normalization parameters are updated**. d is the inner dimensionality of the feed-forward network, which changes the number of parameters of the model.
- **Gradual Unfreezing**² – more parameters finetuned over time. The “unfreezing” starts at the top layer.

Fine-tuning method	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ All parameters	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Adapter layers, $d = 32$	80.52	15.08	79.32	60.40	13.84	17.88	15.54
Adapter layers, $d = 128$	81.51	16.62	79.47	63.03	19.83	27.50	22.63
Adapter layers, $d = 512$	81.54	17.78	79.18	64.30	23.45	33.98	25.81
Adapter layers, $d = 2048$	81.51	16.62	79.47	63.03	19.83	27.50	22.63
Gradual unfreezing	82.50	18.95	79.17	70.79	26.71	39.02	26.93

¹Houlsby et al. “Parameter-efficient transfer learning for NLP”

²Universal language model fine-tuning for text classification

Training strategies

- **Multi-task learning** – train the model at multiple tasks at a time. In T5 this only corresponds to mixing datasets.
- Multiple ways of mixing the datasets:
 - **Proportional mixing** – proportional to the size of each dataset but using an artificial limit K on the size of each dataset.
 - **Temperature-scaled mixing** – using a “temperature” T to control the mixing rate
 - **Equal mixing** – equal probability to every dataset (probably not a good idea)

Mixing strategy	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Baseline (pre-train/fine-tune)	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Equal	76.13	19.02	76.51	63.37	23.89	34.31	26.78
Examples-proportional, $K = 2^{16}$	80.45	19.04	77.25	69.95	24.35	34.99	27.10
Examples-proportional, $K = 2^{17}$	81.56	19.12	77.00	67.91	24.36	35.00	27.25
Examples-proportional, $K = 2^{18}$	81.67	19.07	78.17	67.94	24.57	35.19	27.39
Examples-proportional, $K = 2^{19}$	81.42	19.24	79.78	67.30	25.21	36.30	27.76
Examples-proportional, $K = 2^{20}$	80.80	19.24	80.36	67.38	25.66	36.93	27.68
Examples-proportional, $K = 2^{21}$	79.83	18.79	79.50	65.10	25.82	37.22	27.13
Temperature-scaled, $T = 2$	81.90	19.28	79.42	69.92	25.42	36.72	27.20
Temperature-scaled, $T = 4$	80.56	19.22	77.99	69.54	25.04	35.82	27.45
Temperature-scaled, $T = 8$	77.21	19.10	77.14	66.07	24.55	35.35	27.17

Training strategies

- **Combining multi-task learning with finetuning** – model is pre-trained on all tasks but only fine-tuned on some tasks:
 - Pre-train the model on the mixture of datasets and fine-tuning it on each task
 - Pre-train the model on the mixture of datasets and “leave one out” to be fine-tuned on that task.
 - Pre-training on all datasets

Training strategy	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
★ Unsupervised pre-training + fine-tuning	83.28	19.24	80.88	71.36	26.98	39.82	27.65
Multi-task training	81.42	19.24	79.78	67.30	25.21	36.30	27.76
Multi-task pre-training + fine-tuning	83.11	19.12	80.26	71.03	27.08	39.80	28.07
Leave-one-out multi-task training	81.98	19.05	79.97	71.68	26.93	39.79	27.87
Supervised multi-task pre-training	79.93	18.96	77.38	65.36	26.81	40.13	28.04

- Finetuning after pre-training is comparable to the baseline.
- Supervised pre-training performs worse except for the translation tasks

Scaling

- Possibility to scale in various ways:
 - Using a larger model
 - Increasing batch size
 - Training for more steps
 - Ensembling (combination of models)

Scaling strategy	GLUE	CNNDM	SQuAD	SGLUE	EnDe	EnFr	EnRo
Baseline	83.28	19.24	80.88	71.36	26.98	39.82	27.65
1× size, 4× training steps	85.33	19.33	82.45	74.72	27.08	40.66	27.93
1× size, 4× batch size	84.60	19.42	82.52	74.64	27.07	40.60	27.84
2× size, 2× training steps	86.18	19.66	84.18	77.18	27.52	41.03	28.19
4× size, 1× training steps	85.91	19.73	83.86	78.04	27.47	40.71	28.10
4× ensembled	84.77	20.10	83.09	71.74	28.05	40.53	28.57
4× ensembled, fine-tune only	84.05	19.57	82.36	71.55	27.55	40.22	28.09

- Increasing both model size and training steps improves the results.
- Increasing only batch size or training steps are equally beneficial.
- Ensembling proves to be a way of improving performance without increasing model size or training time.

Scaling – Model Size

Model Name	# Parameters	Hidden Size	# Attention Heads	# Total Layers
Small	60 Million	512	8	12
Base	220 Million	768	12	24
Large	770 Million	1024	16	24
3B	2.8 Billion	1024	32	48
11B	11 Billion	1024	128	48

11 B = 11.000.000.000 !!!!

Conclusion

- **Text-to-text** – simple and easily understandable format, that can be adapted to various tasks.
- **Architectures** – the best architecture for a text-to-text format is the encoder-decoder, that although having twice as many parameters as “encoder-only” (BERT) it has a similar computational cost.
- **Unsupervised objectives** – denoising objectives performed better. Changes in the typical algorithm can provide more efficient training.
- **Datasets** – performance degrades when a dataset is repeated many times during pre-training.
- **Training strategies** – updating all weights performed best. But there are methods that can perform similarly in terms of results but faster during training, thanks to only updating part of the parameters.
- **Scaling** – various ways of scaling up to improve performance.

Research Opportunities

- **Model are very large** – invest in ways of reducing the size of the existing models or create new cheaper models.
- **More efficient knowledge extraction** – new pre-training objectives that can leverage the text in a more efficient way, not being necessary to use such a large amount of data.
- **Measure similarity** between pre-training and downstream tasks.
- **Language Agnostic models** – develop models that achieve good performance regardless on the text's language.

A Simple Coding Example

- Following this [tutorial](#), we can use Google's free TPU's available in Colab to train a T5 model for QA.
- Basic pipeline for using T5:
 - Use Colab and a Google Cloud Storage Account (300\$ free credit)
 - Install the T5 library
 - Pre-process the data
 - Create the necessary datasets (training, validation, test)
 - Create a task or mixture
 - Define model
 - Finetune on the task
 - Save model
 - Evaluate
- You can also play a trivia question game with the model [here](#).

Thank You!