

Weight Uncertainty in Neural Networks

Nova Search Reading Group

by: Simão Gonçalves
date: 2 July 2020

Motivation

- Neural Networks lack a confidence metric on their predictions
 - Critical applications need this: self driving, fraud detection, credit scoring...

Some background first

- Marginal probability:

$$\forall x \in \mathbf{x}, P(\mathbf{x} = x) = \sum_y P(\mathbf{x} = x, y = y)$$


Joint distribution: $P(H, L)$

H \ L	Red	Yellow	Green	Marginal probability $P(H)$
Not Hit	0.198	0.09	0.14	0.428
Hit	0.002	0.01	0.56	0.572
Total	0.2	0.1	0.7	1

Some background first

- Conditional probability:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$



What portion of B contains A?

Some background first

- Bayes theorem:

$$P(A|B) = \frac{P(B|A) P(A)}{P(B)}$$

Some background first

- Bayes theorem:

The diagram shows the formula for Bayes' Theorem:
$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$
 Each part of the formula is labeled with an orange arrow pointing to it. The label for $P(H|E)$ is 'Posterior Probability of 'H' given the evidence'. The label for $P(H)$ is 'Prior Probability'. The label for $P(E|H)$ is 'Likelihood of the evidence 'E' if the Hypothesis 'H' is true'. The label for $P(E)$ is 'Prior probability that the evidence itself is true'.

Prior Probability

Likelihood of the evidence 'E' if the Hypothesis 'H' is true

$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$

Posterior Probability of 'H' given the evidence

Prior probability that the evidence itself is true

Some background first

- Bayes theorem:

Example with coin tosses

The diagram shows the formula for Bayes' Theorem:
$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$
 Each part of the formula is labeled with an orange arrow pointing to it. The label for $P(H|E)$ is 'Posterior Probability of 'H' given the evidence'. The label for $P(H)$ is 'Prior Probability'. The label for $P(E|H)$ is 'Likelihood of the evidence 'E' if the Hypothesis 'H' is true'. The label for $P(E)$ is 'Prior probability that the evidence itself is true'.

Prior Probability

Likelihood of the evidence 'E' if the Hypothesis 'H' is true

$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$

Posterior Probability of 'H' given the evidence

Prior probability that the evidence itself is true

Some background first

Example with coin tosses

- Bayes theorem:

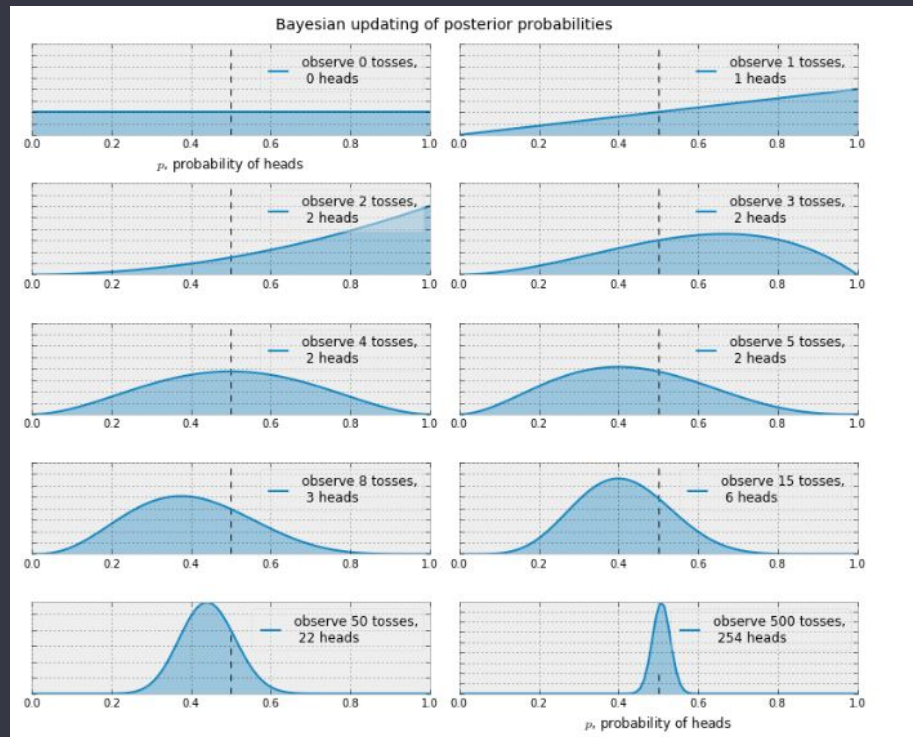
$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$

Prior Probability


Likelihood of the evidence 'E' if the Hypothesis 'H' is true

Posterior Probability of 'H' given the evidence

Priori probability that the evidence itself is true



Probabilistic View on Neural Networks



$$p_z(y|x)$$

Probabilistic View on Neural Networks

$$\begin{aligned} z^{MLE} &= \underset{z}{\operatorname{arg\,max}} \log P(D|z) \\ &= \underset{z}{\operatorname{arg\,max}} \sum_i \log P(y_i|x_i, z) \end{aligned}$$

Probabilistic View on Neural Networks

Bayesian setup

$$\{x^{(i)}, y^{(i)}\}_{i=1:M} \longrightarrow p(y|x)$$

Posterior

$$p(z|D) = \frac{p(D|z)p(z)}{p(D)}$$

Predictions

$$\begin{aligned} p(\hat{y}|\hat{x}, \mathcal{D}) &= \int_{\mathcal{Z}} p(\hat{y}|\hat{x}, z) p(z|\mathcal{D}) dz \\ &= E_{p(z|\mathcal{D})} [P(\hat{y}|\hat{x}, z)] \end{aligned}$$

- Ensemble of predictions

Bad news: it is intractable.

Variational Learning




Fig. 1: Intuition of variational learning.

Source: <https://medium.com/neuralspace/probabilistic-deep-learning-bayes-by-backprop-c4a3de0d9743>

Variational Learning

How do we find q ?

1. Minimize KL divergence between the two distributions.

$$\theta_{\text{opt}} = \underset{\theta \in \Theta}{\operatorname{argmin}} \operatorname{KL}(q_{\theta}(z) || p(z|x))$$

Variational Learning

$$\begin{aligned}\text{KL}\left(q_{\theta}(z)||p(z|x)\right) &= \int q_{\theta}(z)\log\frac{q_{\theta}(z)}{p(z|x)}dz \\ &= E_{q_{\theta}(z)}\left[\log\frac{q_{\theta}(z)}{p(z|x)}\right]\end{aligned}$$

Variational Learning

$$\begin{aligned}\text{KL}\left(q_{\theta}(z)||p(z|x)\right) &= \int q_{\theta}(z)\log\frac{q_{\theta}(z)}{p(z|x)}dz \\ &= E_{q_{\theta}(z)}\left[\log\frac{q_{\theta}(z)}{p(z|x)}\right]\end{aligned}$$

• • •

$$\begin{aligned}&= -\left(E_{q_{\theta}(z)}\left[\log p(z, x)\right] - E_{q_{\theta}(z)}\left[\log q_{\theta}(z)\right]\right) \\ &+ \log p(x)\end{aligned}$$

Variational Learning

Evidence Lower Bound (ELBO): $E_{q_{\theta}(z)} [\log p(z, x)] - E_{q_{\theta}(z)} [\log q_{\theta}(z)]$

$$= E_{q_{\theta}(x|z)} [p(x|z)] + KL[p(x)||q(z)]$$

Variational Learning

Evidence Lower Bound (ELBO): $E_{q_{\theta}(z)} [\log p(z, x)] - E_{q_{\theta}(z)} [\log q_{\theta}(z)]$

$$= E_{q_{\theta}(x|z)} [p(x|z)] + KL[p(x)||q(z)]$$

Maximizing ELBO minimizes the KL divergence!

Cost Function

As we saw earlier:

$$\begin{aligned}\theta_{\text{opt}} &= \underset{\theta}{\operatorname{argmin}} \operatorname{KL}(q_{\theta}(z) || p(z|\mathcal{D})) \\ &= \underset{\theta}{\operatorname{argmin}} \int q_{\theta}(z) \log \frac{q_{\theta}(z)}{p(z)p(\mathcal{D}|z)} dz \\ &= \underset{\theta}{\operatorname{argmin}} \operatorname{KL}[q_{\theta}(z) || p(z)] - E_{q_{\theta}(z)}[\log p(\mathcal{D}|z)]\end{aligned}$$

Cost function

$$F(\mathcal{D}, \theta) = \operatorname{KL}[q_{\theta}(z) || p(z)] - E_{q_{\theta}(z)}[\log p(\mathcal{D}|z)]$$

Gradient Learning

$$\frac{\partial}{\partial z} F(D, \theta)$$



We can't compute gradients of expectations.

Gradient Learning

Authors' contribution: Bayes by Backprop

Reparameterize with a deterministic function:

$$\begin{aligned}\epsilon &\sim q(\epsilon) \\ w &= t(\theta, \epsilon) \text{ - (deterministic)}\end{aligned}$$

Under certain conditions:

$$\frac{\partial}{\partial \theta} \mathbb{E}_{q(\mathbf{w}|\theta)}[f(\mathbf{w}, \theta)] = \mathbb{E}_{q(\epsilon)} \left[\frac{\partial f(\mathbf{w}, \theta)}{\partial \mathbf{w}} \frac{\partial \mathbf{w}}{\partial \theta} + \frac{\partial f(\mathbf{w}, \theta)}{\partial \theta} \right]$$

Gradient Learning

Authors' contribution: Bayes by Backprop

Let's say our weights follow a normal distribution:

$$\mathbf{w} \sim N(\mu, \sigma^2)$$

$$\epsilon \sim N(0, I)$$

$$f(\epsilon) = \mathbf{w} = \mu + \sigma \cdot \epsilon$$

Gradient Learning


Authors' contribution: Bayes by Backprop

Let's say our weights follow a normal distribution:

$$\mathbf{w} \sim N(\mu, \sigma^2)$$

$$\epsilon \sim N(0, I)$$

$$f(\epsilon) = \mathbf{w} = \mu + \sigma \cdot \epsilon$$



$$\mathbf{w} = \mu + \log(1 + \exp(\rho)) \circ \epsilon.$$

Gradient Learning


Authors' contribution: Bayes by Backprop


Let's say our weights follow a normal distribution:

$$\mathbf{w} \sim N(\mu, \sigma^2)$$

$$\epsilon \sim N(0, I)$$

$$f(\epsilon) = \mathbf{w} = \mu + \sigma \cdot \epsilon$$


$$\sigma = \log(1 + \exp(\rho))$$

$$\mathbf{w} = \mu + \log(1 + \exp(\rho)) \circ \epsilon.$$


$$\Delta_{\mu} = \frac{\partial f(\mathbf{w}, \theta)}{\partial \mathbf{w}} + \frac{\partial f(\mathbf{w}, \theta)}{\partial \mu}.$$

$$\Delta_{\rho} = \frac{\partial f(\mathbf{w}, \theta)}{\partial \mathbf{w}} \frac{\epsilon}{1 + \exp(-\rho)} + \frac{\partial f(\mathbf{w}, \theta)}{\partial \rho}$$

Gradient Learning


Authors' contribution: Bayes by Backprop


Let's say our weights follow a normal distribution:

$$\mathbf{w} \sim N(\mu, \sigma^2)$$

$$\epsilon \sim N(0, I)$$

$$f(\epsilon) = \mathbf{w} = \mu + \sigma \cdot \epsilon$$


$$\sigma = \log(1 + \exp(\rho))$$

$$\mathbf{w} = \mu + \log(1 + \exp(\rho)) \circ \epsilon.$$


$$\Delta_{\mu} = \frac{\partial f(\mathbf{w}, \theta)}{\partial \mathbf{w}} + \frac{\partial f(\mathbf{w}, \theta)}{\partial \mu}.$$

$$\Delta_{\rho} = \frac{\partial f(\mathbf{w}, \theta)}{\partial \mathbf{w}} \frac{\epsilon}{1 + \exp(-\rho)} + \frac{\partial f(\mathbf{w}, \theta)}{\partial \rho}$$

$$\mu \leftarrow \mu - \alpha \Delta_{\mu}$$

$$\rho \leftarrow \rho - \alpha \Delta_{\rho}$$

Going back to the cost function

$$F(\mathcal{D}, \theta) = \text{KL}[q_{\theta}(z)||p(z)] - E_{q_{\theta}(z)}[\log p(\mathcal{D}|z)]$$

$$= E[\log q_{\theta}(z^{(i)}|D)] - E[\log p(z^{(i)})] - E[\log p(D|z^{(i)})]$$

We approximate it by taking samples:

$$F(\mathcal{D}, \theta) \approx f(z, \theta) = \sum_{i=1}^n \log q_{\theta}(z^{(i)}) - \log p(z^{(i)}) - \log p(\mathcal{D}|z^{(i)})$$

How do we build this?

```
class Linear_BBB(nn.Module):
    """
    Layer of our BNN.
    """
    def __init__(self, input_features, output_features, prior_var=1.):
        """
        Initialization of our layer : our prior is a normal distribution
        centered in 0 and of variance 20.
        """
        # initialize layers
        super().__init__()
        # set input and output dimensions
        self.input_features = input_features
        self.output_features = output_features

        # initialize mu and rho parameters for the weights of the layer
        self.w_mu = nn.Parameter(torch.zeros(output_features, input_features))
        self.w_rho = nn.Parameter(torch.zeros(output_features, input_features))

        # initialize mu and rho parameters for the layer's bias
        self.b_mu = nn.Parameter(torch.zeros(output_features))
        self.b_rho = nn.Parameter(torch.zeros(output_features))

        # initialize weight samples (these will be calculated whenever the layer makes a prediction)
        self.w = None
        self.b = None

        # initialize prior distribution for all of the weights and biases
        self.prior = torch.distributions.Normal(0, prior_var)
```

How do we build this?

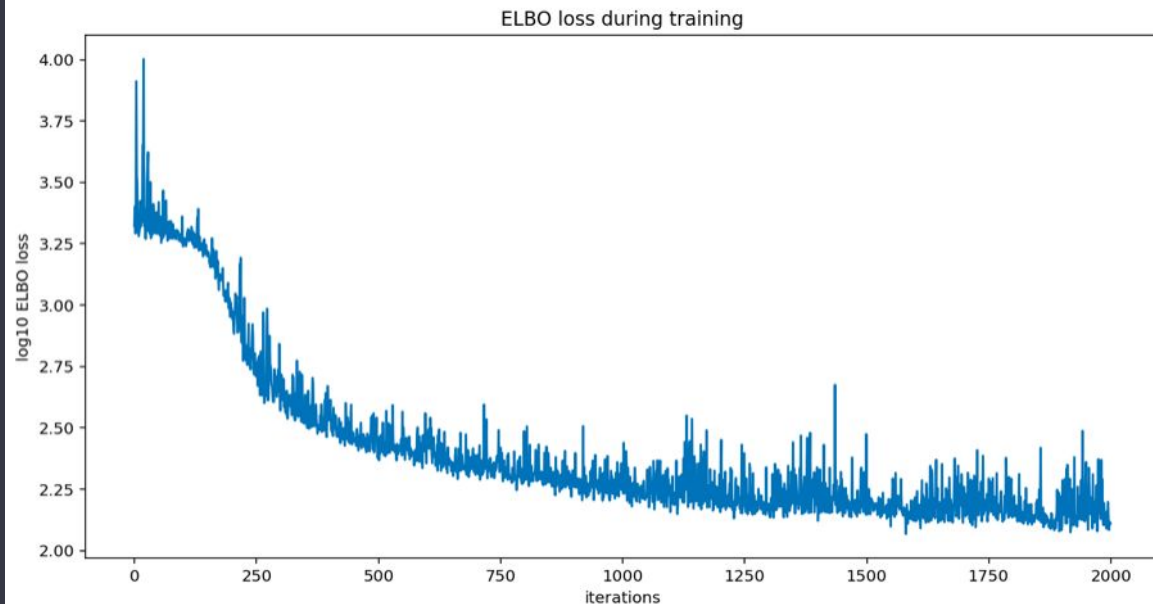
```
def forward(self, input):  
    """  
    Optimization process  
    """  
    # sample weights  
    w_epsilon = Normal(0,1).sample(self.w_mu.shape)  
    self.w = self.w_mu + torch.log(1+torch.exp(self.w_rho)) * w_epsilon  
  
    # sample bias  
    b_epsilon = Normal(0,1).sample(self.b_mu.shape)  
    self.b = self.b_mu + torch.log(1+torch.exp(self.b_rho)) * b_epsilon  
  
    # record log prior by evaluating log pdf of prior at sampled weight and bias  
    w_log_prior = self.prior.log_prob(self.w)  
    b_log_prior = self.prior.log_prob(self.b)  
    self.log_prior = torch.sum(w_log_prior) + torch.sum(b_log_prior)  
  
    # record log variational posterior by evaluating log pdf of normal distribution  
    # defined by parameters with respect at the sampled values  
    self.w_post = Normal(self.w_mu.data, torch.log(1+torch.exp(self.w_rho)))  
    self.b_post = Normal(self.b_mu.data, torch.log(1+torch.exp(self.b_rho)))  
    self.log_post = self.w_post.log_prob(self.w).sum() + self.b_post.log_prob(self.b).sum()  
  
    return F.linear(input, self.w, self.b)
```

How do we build this?

```
def sample_elbo(self, input, target, samples):  
    # we calculate the negative elbo, which will be our loss function  
    # initialize tensors  
    outputs = torch.zeros(samples, target.shape[0])  
    log_priors = torch.zeros(samples)  
    log_posts = torch.zeros(samples)  
    log_likes = torch.zeros(samples)  
    # make predictions and calculate prior, posterior, and likelihood for a given number of samples  
    for i in range(samples):  
        outputs[i] = self(input).reshape(-1) # make predictions  
        log_priors[i] = self.log_prior() # get log prior  
        log_posts[i] = self.log_post() # get log variational posterior  
        log_likes[i] = Normal(outputs[i], self.noise_tol).log_prob(target.reshape(-1)).sum() # calculate the log likelihood  
    # calculate monte carlo estimate of prior posterior and likelihood  
    log_prior = log_priors.mean()  
    log_post = log_posts.mean()  
    log_like = log_likes.mean()  
    # calculate the negative elbo (which is our loss function)  
    loss = log_post - log_prior - log_like  
    return loss
```

How do we build this?

```
plt.figure(figsize=(12,6))
plt.plot(range(len(loss_history)), np.log10(loss_history))
plt.title('ELBO loss during training')
plt.xlabel('iterations')
plt.ylabel('log10 ELBO loss')
plt.show()
```



Bayesian deep learning with CNNs - MNIST




Bayesian deep learning with CNNs - MNIST

- What happens if we input an image out of distribution of the training set?

Bayesian deep learning with CNNs - MNIST

- What happens if we input an image out of distribution of the training set?



Undecided.

